

## Q2

November 2, 2023

```
[1]: import numpy as np
import sys
import pdb
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, accuracy_score, f1_score
import matplotlib.pyplot as plt
from sklearn.preprocessing import OneHotEncoder
import random
```

```
[2]: def get_data(x_path, y_path):
    '''
    Args:
        x_path: path to x file
        y_path: path to y file
    Returns:
        x: np array of [NUM_OF_SAMPLES x n]
        y: np array of [NUM_OF_SAMPLES]
    '''
    x = np.load(x_path)
    y = np.load(y_path)

    y = y.astype('float')
    x = x.astype('float')

    #normalize x:
    x = 2*(0.5 - x/255)
    return x, y
```

```
[3]: def get_metric(y_true, y_pred):
    '''
    Args:
        y_true: np array of [NUM_SAMPLES x r] (one hot)
                or np array of [NUM_SAMPLES]
        y_pred: np array of [NUM_SAMPLES x r] (one hot)
                or np array of [NUM_SAMPLES]
    '''
```

```

results = classification_report(y_pred, y_true)
print(results)

```

### Preparing the X\_train, y\_train, X\_test, y\_test

```

[4]: x_train_path = "/home/tkarthikeyan/IIT Delhi/COL774-Machine Learning/Assignment_
      ↪3/Data_b/x_train.npy"
y_train_path = "/home/tkarthikeyan/IIT Delhi/COL774-Machine Learning/Assignment_
      ↪3/Data_b/y_train.npy"

X_train, y_train = get_data(x_train_path, y_train_path)

x_test_path = "/home/tkarthikeyan/IIT Delhi/COL774-Machine Learning/Assignment_
      ↪3/Data_b/x_test.npy"
y_test_path = "/home/tkarthikeyan/IIT Delhi/COL774-Machine Learning/Assignment_
      ↪3/Data_b/y_test.npy"

X_test, y_test = get_data(x_test_path, y_test_path)

#you might need one hot encoded y in part a,b,c,d,e
label_encoder = OneHotEncoder(sparse_output = False)
label_encoder.fit(np.expand_dims(y_train, axis = -1))

y_train_onehot = label_encoder.transform(np.expand_dims(y_train, axis = -1))
y_test_onehot = label_encoder.transform(np.expand_dims(y_test, axis = -1))

```

```

[5]: class NeuralNetwork:
      def __init__(self, n, n_hidden_nodes, r, M):
          #Number of nodes in the architecture
          self.n = n
          self.n_hidden_nodes = n_hidden_nodes
          self.r = r

          #Mini batch size
          self.M = M

          #Weights and biases
          self.W = dict()
          self.b = dict()

      def initialize_weights_and_biases(self):
          n_nodes = [self.n] + self.n_hidden_nodes + [self.r]

          #Initialize weights
          for i in range(1, len(n_nodes)):
              self.W[str(i)] = np.random.uniform(low=-0.1, high=0.1,
          ↪size=(n_nodes[i], n_nodes[i-1]))

```

```

        #Initialize biases
        for i in range(1,len(n_nodes)):
            self.b[str(i)] = np.zeros((n_nodes[i],1))

    @staticmethod
    def sigmoid(x, derivative = False):
        if derivative == False:
            return 1 / (1 + np.exp(-x))
        else:
            return NeuralNetwork.sigmoid(x, derivative = False) * (1 -
↪NeuralNetwork.sigmoid(x, derivative = False))

    @staticmethod
    def relu(x, derivative = False):
        if derivative == True:
            return np.where(x > 0, 1, np.where(x < 0, 0, np.random.
↪random_sample()))
        else:
            return np.where(x <= 0, 0, x)

    @staticmethod
    def softmax(Z):
        return np.exp(Z) / np.sum(np.exp(Z), axis=0)

    def train(self, X_train, y_train, epoch_mode = True, activation="sigmoid",
↪EPOCHS = 200, alpha = 0.01, stopping_threshold = None,
↪adaptive_learning=False, printafter=20):
        self.initialize_weights_and_biases()

        a = dict()
        z = dict()
        del_z = dict()
        del_b = dict()
        del_W = dict()

        if epoch_mode == True:
            for epoch in range(EPOCHS):
                for i in range(0, X_train.shape[0], self.M):
                    y_actual = y_train[i:i+self.M,:].T

                    #Forward
                    a["0"] = X_train[i:i+self.M,:].T

                    for j in range(1,len(self.n_hidden_nodes)+1):
                        z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) +
↪self.b[str(j)]

```

```

        if activation == "relu":
            a[str(j)] = NeuralNetwork.relu(z[str(j)])
        else:
            a[str(j)] = NeuralNetwork.sigmoid(z[str(j)])

        j += 1
        z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) + self.
↪b[str(j)]

        a[str(j)] = NeuralNetwork.softmax(z[str(j)])

        #Backward
        del_z[str(j)] = a[str(j)] - y_actual
        del_b[str(j)] = np.sum(del_z[str(j)], axis = 1).
↪reshape(-1,1)

        del_W[str(j)] = np.matmul(del_z[str(j)], a[str(j-1)].T)

        for k in range(j-1,0,-1):
            if activation == "relu":
                del_z[str(k)] = np.matmul(self.W[str(k+1)].T,
↪del_z[str(k+1)])*(NeuralNetwork.relu(z[str(k)], derivative=True))
            else:
                del_z[str(k)] = np.matmul(self.W[str(k+1)].T,
↪del_z[str(k+1)])*(NeuralNetwork.sigmoid(z[str(k)], derivative=True))
                del_b[str(k)] = np.sum(del_z[str(k)], axis = 1).
↪reshape(-1,1)

                del_W[str(k)] = np.matmul(del_z[str(k)], a[str(k-1)].T)

        #Update
        for l in range(1,len(self.n_hidden_nodes)+2):
            if adaptive_learning == False:
                self.W[str(l)] = self.W[str(l)] - alpha *
↪del_W[str(l)]

                self.b[str(l)] = self.b[str(l)] - alpha *
↪del_b[str(l)]

            else:
                self.W[str(l)] = self.W[str(l)] - (alpha/np.
↪sqrt(epoch)) * del_W[str(l)]
                self.b[str(l)] = self.b[str(l)] - (alpha/np.
↪sqrt(epoch)) * del_b[str(l)]

        y_pred, softmax_output = NN.predict(X_train,
↪activation=activation)
        softmax_loss = NeuralNetwork.
↪compute_softmax_loss(softmax_output, y_train_onehot)

        if epoch%printafter==0:
            print(f"epoch {epoch}")

```

```

        print("accuracy on train data:␣
↪",accuracy_score(y_train_onehot, y_pred))
        if adaptive_learning:
            print("learning rate: ",(alpha/np.sqrt(epoch)))
            print("softmax loss: ",softmax_loss)
            print("\n")

    else:
        epoch = 1
        window = 5
        loss_avg = 0
        while(True):
            for i in range(0, X_train.shape[0], self.M):
                y_actual = y_train[i:i+self.M,:].T

                #Forward
                a["0"] = X_train[i:i+self.M,:].T

                for j in range(1,len(self.n_hidden_nodes)+1):
                    z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) +␣
↪self.b[str(j)]

                    if activation == "relu":
                        a[str(j)] = NeuralNetwork.relu(z[str(j)])
                    else:
                        a[str(j)] = NeuralNetwork.sigmoid(z[str(j)])

                j += 1
                z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) + self.
↪b[str(j)]

                a[str(j)] = NeuralNetwork.softmax(z[str(j)])

                #Backward
                del_z[str(j)] = a[str(j)] - y_actual
                del_b[str(j)] = np.sum(del_z[str(j)], axis = 1).
↪reshape(-1,1)

                del_W[str(j)] = np.matmul(del_z[str(j)], a[str(j-1)].T)

                for k in range(j-1,0,-1):
                    if activation == "relu":
                        del_z[str(k)] = np.matmul(self.W[str(k+1)].T,␣
↪del_z[str(k+1)])*(NeuralNetwork.relu(z[str(k)], derivative=True))
                    else:
                        del_z[str(k)] = np.matmul(self.W[str(k+1)].T,␣
↪del_z[str(k+1)])*(NeuralNetwork.sigmoid(z[str(k)], derivative=True))
                        del_b[str(k)] = np.sum(del_z[str(k)], axis = 1).
↪reshape(-1,1)

```

```

        del_W[str(k)] = np.matmul(del_z[str(k)], a[str(k-1)].T)

        #Update
        for l in range(1,len(self.n_hidden_nodes)+2):
            if adaptive_learning == False:
                self.W[str(l)] = self.W[str(l)] - alpha *
→del_W[str(l)]
                self.b[str(l)] = self.b[str(l)] - alpha *
→del_b[str(l)]
            else:
                self.W[str(l)] = self.W[str(l)] - (alpha/np.
→sqrt(epoch)) * del_W[str(l)]
                self.b[str(l)] = self.b[str(l)] - (alpha/np.
→sqrt(epoch)) * del_b[str(l)]

        y_pred, softmax_output = NN.predict(X_train,
→activation=activation)
        softmax_loss = NeuralNetwork.
→compute_softmax_loss(softmax_output, y_train_onehot)
        if epoch%printafter==0:
            print(f"epoch {epoch}")
            print("accuracy on train data:
→",accuracy_score(y_train_onehot, y_pred))
            if adaptive_learning:
                print("learning rate: ",(alpha/np.sqrt(epoch)))
                print("softmax loss: ",softmax_loss)
                print("\n")

        if epoch <= window:
            loss_avg += softmax_loss
            if epoch == window:
                loss_avg /= window

        #End the training
        if epoch > window:
            new_loss_avg = ((window - 1)*loss_avg + softmax_loss)/window
            diff_avg_loss = abs(new_loss_avg - loss_avg)
            #print("diff avg loss:",diff_avg_loss)
            if diff_avg_loss < stopping_threshold or epoch > EPOCHS:
                print("Convergence criteria satisfied!")
                print(f"epoch {epoch}")
                print("accuracy on train data:
→",accuracy_score(y_train_onehot, y_pred))
                if adaptive_learning:
                    print("learning rate: ",(alpha/np.sqrt(epoch)))
                    print("softmax loss: ",softmax_loss)

```

```

        print("\n")
        break

    epoch += 1

def predict(self, X_test, activation="sigmoid"):
    y_pred = np.zeros((X_test.shape[0], self.r))
    softmax_output = np.zeros((X_test.shape[0], self.r))
    z = dict()
    a = dict()

    for i in range(X_test.shape[0]):
        a[str(0)] = X_test[i:i+1,:].T

        for j in range(1, len(self.n_hidden_nodes)+1):
            z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) + self.
→b[str(j)]

            if activation == "relu":
                a[str(j)] = NeuralNetwork.relu(z[str(j)])
            else:
                a[str(j)] = NeuralNetwork.sigmoid(z[str(j)])

        j += 1
        z[str(j)] = np.matmul(self.W[str(j)], a[str(j-1)]) + self.b[str(j)]
        a[str(j)] = NeuralNetwork.softmax(z[str(j)])

        softmax_output[i] = a[str(j)].flatten()
        y_pred[i][np.argmax(a[str(j)])] = 1

    return y_pred, softmax_output

@staticmethod
def compute_softmax_loss(softmax_output, y_pred):
    #Softmax loss
    sm_loss = 0
    for i in range(y_pred.shape[0]):
        sm_loss = -1*np.log2(softmax_output[i][np.argmax(y_pred[i])])

    return sm_loss/(y_pred.shape[0])

```

```

[28]: NN = NeuralNetwork(n = 1024, n_hidden_nodes = [100,50] , r = 5, M = 32)
NN.train(X_train, y_train_onehot, activation="sigmoid", EPOCHS=100, alpha=0.
→001, printafter=10)
y_pred, softmax_output = NN.predict(X_test, activation="sigmoid")
print("accuracy on test data: ", accuracy_score(y_test_onehot, y_pred))
get_metric(y_test_onehot, y_pred)

```

epoch 0  
accuracy on train data: 0.2091  
softmax loss: 0.00023525478839202835

epoch 10  
accuracy on train data: 0.5923  
softmax loss: 0.00011264389363753989

epoch 20  
accuracy on train data: 0.6572  
softmax loss: 5.831073565883006e-05

epoch 30  
accuracy on train data: 0.6702  
softmax loss: 4.656056662492289e-05

epoch 40  
accuracy on train data: 0.6758  
softmax loss: 3.9645451044889796e-05

epoch 50  
accuracy on train data: 0.6803  
softmax loss: 3.4488996828750714e-05

epoch 60  
accuracy on train data: 0.6869  
softmax loss: 3.0593861550354343e-05

epoch 70  
accuracy on train data: 0.6908  
softmax loss: 2.8066258449785907e-05

epoch 80  
accuracy on train data: 0.6976  
softmax loss: 2.7373041902283903e-05

epoch 90  
accuracy on train data: 0.7037  
softmax loss: 2.8412086233645623e-05



accuracy on test data: 0.714

	precision	recall	f1-score	support
0	0.93	0.95	0.94	224
1	0.75	0.76	0.76	196
2	0.56	0.64	0.60	176
3	0.49	0.52	0.51	177
4	0.80	0.66	0.72	227
micro avg	0.71	0.71	0.71	1000
macro avg	0.71	0.70	0.70	1000
weighted avg	0.72	0.71	0.72	1000
samples avg	0.71	0.71	0.71	1000

```
[34]: print("f1_score: ",f1_score(y_test_onehot, y_pred, average="macro"))
```

f1\_score: 0.7029924239467376

### Experimenting with single hidden layer

```
[41]: hidden_layers = [[1],[5],[10],[50],[100]]
number_of_hidden_units = [1,5,10,50,100]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    NN = NeuralNetwork(n = 1024, n_hidden_nodes = hidden_layer , r = 5, M = 32)
    NN.train(X_train, y_train_onehot, epoch_mode= False, activation="sigmoid",
    ↪alpha = 0.01, stopping_threshold = 1.0e-06, printafter=50)
    y_pred_train, _ = NN.predict(X_train)
    y_pred_test, _ = NN.predict(X_test)

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
    get_metric(y_test_onehot, y_pred_test)
    f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))

    print("\n")
```

Hidden layer: [1]

epoch 50

accuracy on train data: 0.4706

softmax loss: 0.00014958531153011078

epoch 100

accuracy on train data: 0.2091

softmax loss: 0.00023287816358262217

epoch 150

accuracy on train data: 0.2091

softmax loss: 0.0002328781641601466

epoch 200

accuracy on train data: 0.2091

softmax loss: 0.00023287816494873916

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.2091

softmax loss: 0.00023287816496746151

accuracy on train data: 0.2091

metrics for train data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.21	0.35	10000
micro avg	0.21	0.21	0.21	10000
macro avg	0.20	0.04	0.07	10000
weighted avg	1.00	0.21	0.35	10000
samples avg	0.21	0.21	0.21	10000

accuracy on test data: 0.187

metrics for test data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0

2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.19	0.32	1000
micro avg	0.19	0.19	0.19	1000
macro avg	0.20	0.04	0.06	1000
weighted avg	1.00	0.19	0.32	1000
samples avg	0.19	0.19	0.19	1000

Hidden layer: [5]

```
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

```
epoch 50
accuracy on train data: 0.6775
softmax loss: 5.1494534443521053e-05
```

```
epoch 100
accuracy on train data: 0.6858
softmax loss: 3.220427274554847e-05
```

```
epoch 150
accuracy on train data: 0.6952
softmax loss: 2.7774495226855474e-05
```

```
epoch 200
accuracy on train data: 0.6982
softmax loss: 2.603425353508207e-05
```

Convergence criteria satisfied!

```
epoch 201
accuracy on train data: 0.6977
softmax loss: 2.6049241749681784e-05
```

accuracy on train data: 0.6977

metrics for train data:

	precision	recall	f1-score	support
0	0.95	0.86	0.90	2176
1	0.64	0.77	0.70	1652
2	0.73	0.56	0.63	2544
3	0.41	0.56	0.47	1447
4	0.77	0.73	0.75	2181
micro avg	0.70	0.70	0.70	10000
macro avg	0.70	0.70	0.69	10000
weighted avg	0.72	0.70	0.70	10000
samples avg	0.70	0.70	0.70	10000

accuracy on test data: 0.689

metrics for test data:

	precision	recall	f1-score	support
0	0.94	0.87	0.91	246
1	0.63	0.76	0.69	163
2	0.72	0.58	0.64	248
3	0.41	0.49	0.45	158
4	0.69	0.70	0.69	185
micro avg	0.69	0.69	0.69	1000
macro avg	0.68	0.68	0.68	1000
weighted avg	0.71	0.69	0.69	1000
samples avg	0.69	0.69	0.69	1000

Hidden layer: [10]

epoch 50

accuracy on train data: 0.6885

softmax loss: 2.3303944233466556e-05

epoch 100

accuracy on train data: 0.7189

softmax loss: 2.3564423526932722e-05

epoch 150

accuracy on train data: 0.7241

softmax loss: 2.151708934034692e-05

epoch 200  
accuracy on train data: 0.7438  
softmax loss: 1.7535741967859166e-05

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.7436  
softmax loss: 1.7259669594328303e-05

accuracy on train data: 0.7436  
metrics for train data:

	precision	recall	f1-score	support
0	0.93	0.95	0.94	1930
1	0.75	0.81	0.78	1852
2	0.67	0.66	0.66	2000
3	0.51	0.59	0.55	1752
4	0.85	0.72	0.78	2466
micro avg	0.74	0.74	0.74	10000
macro avg	0.74	0.74	0.74	10000
weighted avg	0.75	0.74	0.75	10000
samples avg	0.74	0.74	0.74	10000

accuracy on test data: 0.724  
metrics for test data:

	precision	recall	f1-score	support
0	0.89	0.96	0.92	212
1	0.73	0.76	0.75	191
2	0.65	0.64	0.65	204
3	0.51	0.55	0.53	175
4	0.80	0.69	0.74	218
micro avg	0.72	0.72	0.72	1000
macro avg	0.72	0.72	0.72	1000
weighted avg	0.73	0.72	0.72	1000
samples avg	0.72	0.72	0.72	1000

Hidden layer: [50]  
epoch 50  
accuracy on train data: 0.7392  
softmax loss: 2.6987471937410535e-05

epoch 100  
accuracy on train data: 0.7839  
softmax loss: 1.7227472892700306e-05

epoch 150  
accuracy on train data: 0.7706  
softmax loss: 1.6560629492585795e-05

epoch 200  
accuracy on train data: 0.7877  
softmax loss: 8.237749203355322e-06

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.788  
softmax loss: 7.399416514018198e-06

accuracy on train data: 0.788  
metrics for train data:

	precision	recall	f1-score	support
0	0.94	0.98	0.96	1878
1	0.79	0.88	0.84	1783
2	0.64	0.72	0.68	1734
3	0.65	0.62	0.64	2106
4	0.91	0.76	0.83	2499
micro avg	0.79	0.79	0.79	10000
macro avg	0.79	0.79	0.79	10000
weighted avg	0.79	0.79	0.79	10000
samples avg	0.79	0.79	0.79	10000

accuracy on test data: 0.735  
metrics for test data:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	215
1	0.76	0.82	0.79	183
2	0.55	0.65	0.60	168
3	0.63	0.52	0.57	229
4	0.76	0.70	0.73	205

micro avg	0.73	0.73	0.73	1000
macro avg	0.73	0.74	0.73	1000
weighted avg	0.73	0.73	0.73	1000
samples avg	0.73	0.73	0.73	1000

Hidden layer: [100]

epoch 50

accuracy on train data: 0.733

softmax loss: 4.538836493085465e-05

epoch 100

accuracy on train data: 0.7818

softmax loss: 1.8238639023683118e-05

epoch 150

accuracy on train data: 0.8035

softmax loss: 9.2365944713612e-06

epoch 200

accuracy on train data: 0.8139

softmax loss: 9.666839084236848e-06

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.8133

softmax loss: 9.72563184360032e-06

accuracy on train data: 0.8133

metrics for train data:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	2010
1	0.82	0.92	0.87	1766
2	0.74	0.77	0.75	1883
3	0.64	0.67	0.65	1914
4	0.88	0.76	0.82	2427
micro avg	0.81	0.81	0.81	10000
macro avg	0.81	0.82	0.81	10000
weighted avg	0.82	0.81	0.81	10000
samples avg	0.81	0.81	0.81	10000

accuracy on test data: 0.77


metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.97	0.97	228
1	0.74	0.84	0.79	174
2	0.65	0.68	0.67	191
3	0.63	0.59	0.61	197
4	0.82	0.73	0.78	210
micro avg	0.77	0.77	0.77	1000
macro avg	0.76	0.77	0.76	1000
weighted avg	0.77	0.77	0.77	1000
samples avg	0.77	0.77	0.77	1000

```
[43]: #Plot avg f1_scores for different number of hidden units
plt.figure(figsize=(10,5))
plt.plot(number_of_hidden_units, f1_score_train, marker='o', markersize=6,
        color='blue', label='Avg f1_score train data')
plt.plot(number_of_hidden_units, f1_score_test, marker='o', markersize=6,
        color='red', label='Avg f1_score test data')

plt.title('Avg f1_score vs number of hidden units')
plt.xlabel('number of hidden units')
plt.ylabel('Avg f1_score')
plt.xticks(number_of_hidden_units)
plt.ylim(0,1.1)
plt.legend()
plt.show()
```





### Experimenting with hidden layers

```
[44]: hidden_layers = [[512], [512,256], [512,256,128], [512,256,128,64]]
network_depth = [1,2,3,4]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    NN = NeuralNetwork(n = 1024, n_hidden_nodes = hidden_layer , r = 5, M = 32)
    NN.train(X_train, y_train_onehot, epoch_mode=False, activation="sigmoid",
    ↪alpha = 0.01, stopping_threshold = 1.0e-06, printafter=50)
    y_pred_train, _ = NN.predict(X_train)
    y_pred_test, _ = NN.predict(X_test)

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
    get_metric(y_test_onehot, y_pred_test)
    f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))

    print("\n")
```

Hidden layer: [512]

epoch 50

accuracy on train data: 0.7327

softmax loss: 3.9801766116042114e-05

epoch 100

accuracy on train data: 0.7921

softmax loss: 5.780218764128989e-06

epoch 150

accuracy on train data: 0.8118

softmax loss: 3.7589816123345803e-06

epoch 200

accuracy on train data: 0.8299

softmax loss: 1.2992431023633802e-06

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.831

softmax loss: 1.239104246690159e-06

accuracy on train data: 0.831

metrics for train data:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1933
1	0.87	0.92	0.89	1863
2	0.70	0.82	0.75	1669
3	0.75	0.66	0.70	2286
4	0.87	0.81	0.84	2249
micro avg	0.83	0.83	0.83	10000
macro avg	0.83	0.84	0.83	10000
weighted avg	0.83	0.83	0.83	10000
samples avg	0.83	0.83	0.83	10000

accuracy on test data: 0.782

metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	223
1	0.81	0.88	0.84	181

2	0.58	0.72	0.64	162
3	0.73	0.56	0.64	241
4	0.80	0.77	0.78	193
micro avg	0.78	0.78	0.78	1000
macro avg	0.78	0.79	0.78	1000
weighted avg	0.78	0.78	0.78	1000
samples avg	0.78	0.78	0.78	1000

Hidden layer: [512, 256]

epoch 50

accuracy on train data: 0.7299

softmax loss: 2.5517614437811225e-05

epoch 100

accuracy on train data: 0.7836

softmax loss: 6.656917643055116e-06

epoch 150

accuracy on train data: 0.8126

softmax loss: 5.07376437203254e-06

epoch 200

accuracy on train data: 0.8213

softmax loss: 6.25927649641769e-06

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.8206

softmax loss: 6.134507568326737e-06

accuracy on train data: 0.8206

metrics for train data:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1943
1	0.91	0.90	0.90	1997
2	0.71	0.83	0.76	1661
3	0.58	0.68	0.63	1733
4	0.93	0.73	0.82	2666

micro avg	0.82	0.82	0.82	10000
macro avg	0.82	0.82	0.82	10000
weighted avg	0.84	0.82	0.82	10000
samples avg	0.82	0.82	0.82	10000

accuracy on test data: 0.78

metrics for test data:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	226
1	0.88	0.84	0.86	206
2	0.59	0.76	0.66	154
3	0.57	0.57	0.57	186
4	0.85	0.70	0.77	228

micro avg	0.78	0.78	0.78	1000
macro avg	0.77	0.77	0.77	1000
weighted avg	0.79	0.78	0.78	1000
samples avg	0.78	0.78	0.78	1000

Hidden layer: [512, 256, 128]

epoch 50

accuracy on train data: 0.7228

softmax loss: 3.62084087252217e-05

epoch 100

accuracy on train data: 0.7799

softmax loss: 1.917585498575037e-05

epoch 150

accuracy on train data: 0.8019

softmax loss: 1.001024824261873e-05

epoch 200

accuracy on train data: 0.8182

softmax loss: 5.357968292509009e-06

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.8138

softmax loss: 6.170754618629236e-06

accuracy on train data: 0.8138

metrics for train data:

	precision	recall	f1-score	support
0	0.99	0.98	0.99	1996
1	0.88	0.94	0.90	1853
2	0.68	0.80	0.74	1651
3	0.58	0.64	0.61	1809
4	0.94	0.73	0.82	2691
micro avg	0.81	0.81	0.81	10000
macro avg	0.81	0.82	0.81	10000
weighted avg	0.83	0.81	0.82	10000
samples avg	0.81	0.81	0.81	10000

accuracy on test data: 0.778

metrics for test data:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	230
1	0.82	0.89	0.85	182
2	0.58	0.72	0.65	160
3	0.60	0.56	0.58	201
4	0.86	0.70	0.77	227
micro avg	0.78	0.78	0.78	1000
macro avg	0.77	0.77	0.77	1000
weighted avg	0.79	0.78	0.78	1000
samples avg	0.78	0.78	0.78	1000

Hidden layer: [512, 256, 128, 64]

Convergence criteria satisfied!

epoch 25

accuracy on train data: 0.2091

softmax loss: 0.0002359463199165115

accuracy on train data: 0.2091

metrics for train data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0

	4	1.00	0.21	0.35	10000
micro avg		0.21	0.21	0.21	10000
macro avg		0.20	0.04	0.07	10000
weighted avg		1.00	0.21	0.35	10000
samples avg		0.21	0.21	0.21	10000

accuracy on test data: 0.187

metrics for test data:

		precision	recall	f1-score	support
	0	0.00	0.00	0.00	0
	1	0.00	0.00	0.00	0
	2	0.00	0.00	0.00	0
	3	0.00	0.00	0.00	0
	4	1.00	0.19	0.32	1000
micro avg		0.19	0.19	0.19	1000
macro avg		0.20	0.04	0.06	1000
weighted avg		1.00	0.19	0.32	1000
samples avg		0.19	0.19	0.19	1000

/home/tkarthikeyan/.local/lib/python3.8/site-packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))


/home/tkarthikeyan/.local/lib/python3.8/site-packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

```
[45]: #Plot avg f1_scores for different depth
plt.figure(figsize=(10,5))
plt.plot(network_depth, f1_score_train, marker='o', markersize=6, color='blue',
↪label='Avg f1_score train data')
plt.plot(network_depth, f1_score_test, marker='o', markersize=6, color='red',
↪label='Avg f1_score test data')

plt.title('Avg f1_score vs network depth')
plt.xlabel('network depth')
plt.ylabel('Avg f1_score')
plt.xticks(network_depth)
plt.ylim(0,1.1)
```

```
plt.legend()
plt.show()
```



### Adaptive learning

```
[46]: hidden_layers = [[512], [512,256], [512,256,128], [512,256,128,64]]
network_depth = [1,2,3,4]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    NN = NeuralNetwork(n = 1024, n_hidden_nodes = hidden_layer , r = 5, M = 32)
    NN.train(X_train, y_train_onehot, epoch_mode=False, activation="sigmoid",
    ↪adaptive_learning = True, alpha = 0.01, stopping_threshold = 5.0e-06,
    ↪printafter=50)
    y_pred_train, _ = NN.predict(X_train)
    y_pred_test, _ = NN.predict(X_test)

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
```

```

get_metric(y_test_onehot, y_pred_test)
f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))

print("\n")

```

Hidden layer: [512]

epoch 50

accuracy on train data: 0.701

learning rate: 0.001414213562373095

softmax loss: 2.4384530256167004e-05

epoch 100

accuracy on train data: 0.7294

learning rate: 0.001

softmax loss: 2.2590405138106635e-05

epoch 150

accuracy on train data: 0.7535

learning rate: 0.0008164965809277261

softmax loss: 2.4823077935786723e-05

epoch 200

accuracy on train data: 0.7701

learning rate: 0.0007071067811865475

softmax loss: 2.5202653816610038e-05

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.7705

learning rate: 0.0007053456158585983

softmax loss: 2.518159706339217e-05

accuracy on train data: 0.7705

metrics for train data:

	precision	recall	f1-score	support
0	0.97	0.94	0.95	2031
1	0.81	0.83	0.82	1937
2	0.65	0.70	0.68	1809
3	0.57	0.62	0.59	1849
4	0.85	0.75	0.80	2374



micro avg	0.77	0.77	0.77	10000
macro avg	0.77	0.77	0.77	10000
weighted avg	0.78	0.77	0.77	10000
samples avg	0.77	0.77	0.77	10000

accuracy on test data: 0.769

metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.96	0.96	232
1	0.80	0.81	0.81	194
2	0.61	0.69	0.65	176
3	0.63	0.59	0.61	199
4	0.80	0.75	0.77	199

micro avg	0.77	0.77	0.77	1000
macro avg	0.76	0.76	0.76	1000
weighted avg	0.77	0.77	0.77	1000
samples avg	0.77	0.77	0.77	1000

Hidden layer: [512, 256]

epoch 50

accuracy on train data: 0.6913

learning rate: 0.001414213562373095

softmax loss: 2.829604797009096e-05

epoch 100

accuracy on train data: 0.7283

learning rate: 0.001

softmax loss: 2.5478154341155774e-05

epoch 150

accuracy on train data: 0.7609

learning rate: 0.0008164965809277261

softmax loss: 3.3394030775105174e-05

epoch 200

accuracy on train data: 0.7788

learning rate: 0.0007071067811865475

softmax loss: 2.511949838152347e-05

Convergence criteria satisfied!

epoch 201  
accuracy on train data: 0.779  
learning rate: 0.0007053456158585983  
softmax loss: 2.4884937458219116e-05

accuracy on train data: 0.779  
metrics for train data:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	1996
1	0.84	0.86	0.85	1951
2	0.66	0.72	0.69	1801
3	0.56	0.61	0.58	1850
4	0.85	0.74	0.79	2402
micro avg	0.78	0.78	0.78	10000
macro avg	0.78	0.78	0.78	10000
weighted avg	0.79	0.78	0.78	10000
samples avg	0.78	0.78	0.78	10000

accuracy on test data: 0.78  
metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.98	0.98	228
1	0.83	0.83	0.83	199
2	0.63	0.71	0.67	177
3	0.62	0.60	0.61	193
4	0.80	0.74	0.77	203
micro avg	0.78	0.78	0.78	1000
macro avg	0.77	0.77	0.77	1000
weighted avg	0.78	0.78	0.78	1000
samples avg	0.78	0.78	0.78	1000

Hidden layer: [512, 256, 128]  
epoch 50  
accuracy on train data: 0.6908  
learning rate: 0.001414213562373095  
softmax loss: 2.5470862546841987e-05

epoch 100  
accuracy on train data: 0.734  
learning rate: 0.001

softmax loss: 3.018135410329875e-05

epoch 150

accuracy on train data: 0.7671

learning rate: 0.0008164965809277261

softmax loss: 3.57361083756996e-05

epoch 200

accuracy on train data: 0.7839

learning rate: 0.0007071067811865475

softmax loss: 1.8145684474911947e-05

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.7838

learning rate: 0.0007053456158585983

softmax loss: 1.784769632453523e-05

accuracy on train data: 0.7838

metrics for train data:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	1991
1	0.85	0.87	0.86	1923
2	0.70	0.72	0.71	1914
3	0.54	0.62	0.58	1762
4	0.85	0.73	0.79	2410
micro avg	0.78	0.78	0.78	10000
macro avg	0.78	0.78	0.78	10000
weighted avg	0.79	0.78	0.79	10000
samples avg	0.78	0.78	0.78	10000

accuracy on test data: 0.783

metrics for test data:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	229
1	0.83	0.85	0.84	194
2	0.68	0.71	0.70	192
3	0.57	0.60	0.58	176
4	0.80	0.72	0.76	209
micro avg	0.78	0.78	0.78	1000

macro avg	0.77	0.77	0.77	1000
weighted avg	0.79	0.78	0.78	1000
samples avg	0.78	0.78	0.78	1000

Hidden layer: [512, 256, 128, 64]

Convergence criteria satisfied!

epoch 6

accuracy on train data: 0.2091

learning rate: 0.004082482904638631

softmax loss: 0.000232068269599483

accuracy on train data: 0.2091

metrics for train data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.21	0.35	10000
micro avg	0.21	0.21	0.21	10000
macro avg	0.20	0.04	0.07	10000
weighted avg	1.00	0.21	0.35	10000
samples avg	0.21	0.21	0.21	10000

accuracy on test data: 0.187

metrics for test data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.19	0.32	1000
micro avg	0.19	0.19	0.19	1000
macro avg	0.20	0.04	0.06	1000
weighted avg	1.00	0.19	0.32	1000
samples avg	0.19	0.19	0.19	1000


/home/tkarthikeyan/.local/lib/python3.8/site-

```
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

```
[47]: #Plot avg f1_scores for different depth
plt.figure(figsize=(10,5))
plt.plot(network_depth, f1_score_train, marker='o', markersize=6, color='blue',
        ↳label='Avg f1_score train data')
plt.plot(network_depth, f1_score_test, marker='o', markersize=6, color='red',
        ↳label='Avg f1_score test data')

plt.title('Avg f1_score vs network depth')
plt.xlabel('network depth')
plt.ylabel('Avg f1_score')
plt.xticks(network_depth)
plt.ylim(0,1.1)
plt.legend()
plt.show()
```



**Relu activation function**

```
[48]: hidden_layers = [[512], [512,256], [512,256,128], [512,256,128,64]]
network_depth = [1,2,3,4]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    NN = NeuralNetwork(n = 1024, n_hidden_nodes = hidden_layer , r = 5, M = 32)
    NN.train(X_train, y_train_onehot, epoch_mode= False, activation="relu",
    ↪adaptive_learning = True, alpha = 0.01, stopping_threshold = 5.0e-06,
    ↪printafter=50)
    y_pred_train, _ = NN.predict(X_train, activation="relu")
    y_pred_test, _ = NN.predict(X_test, activation="relu")

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
    get_metric(y_test_onehot, y_pred_test)
    f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))

    print("\n")
```

```
Hidden layer: [512]
Convergence criteria satisfied!
epoch 6
accuracy on train data: 0.2091
learning rate: 0.004082482904638631
softmax loss: 0.00023208783009197323
```

```
accuracy on train data: 0.2091
metrics for train data:
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.21	0.35	10000
micro avg	0.21	0.21	0.21	10000
macro avg	0.20	0.04	0.07	10000

weighted avg	1.00	0.21	0.35	10000
samples avg	0.21	0.21	0.21	10000

accuracy on test data: 0.187

metrics for test data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.19	0.32	1000
micro avg	0.19	0.19	0.19	1000
macro avg	0.20	0.04	0.06	1000
weighted avg	1.00	0.19	0.32	1000
samples avg	0.19	0.19	0.19	1000

Hidden layer: [512, 256]

/home/tkarthikeyan/.local/lib/python3.8/site-

packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

/home/tkarthikeyan/.local/lib/python3.8/site-

packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

Convergence criteria satisfied!

epoch 6

accuracy on train data: 0.2091

learning rate: 0.004082482904638631

softmax loss: 0.00023044619233218206

accuracy on train data: 0.2091

metrics for train data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.21	0.35	10000

micro avg	0.21	0.21	0.21	10000
macro avg	0.20	0.04	0.07	10000
weighted avg	1.00	0.21	0.35	10000
samples avg	0.21	0.21	0.21	10000

accuracy on test data: 0.187

metrics for test data:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
1	0.00	0.00	0.00	0
2	0.00	0.00	0.00	0
3	0.00	0.00	0.00	0
4	1.00	0.19	0.32	1000

micro avg	0.19	0.19	0.19	1000
macro avg	0.20	0.04	0.06	1000
weighted avg	1.00	0.19	0.32	1000
samples avg	0.19	0.19	0.19	1000

Hidden layer: [512, 256, 128]

/home/tkarthikeyan/.local/lib/python3.8/site-packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

/home/tkarthikeyan/.local/lib/python3.8/site-packages/sklearn/metrics/\_classification.py:1469: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

epoch 50

accuracy on train data: 0.7289

learning rate: 0.001414213562373095

softmax loss: 3.145936113878991e-05

epoch 100

accuracy on train data: 0.7583

learning rate: 0.001

softmax loss: 3.272657747861718e-05

epoch 150



accuracy on train data: 0.7812  
learning rate: 0.0008164965809277261  
softmax loss: 1.1503443467968143e-05

epoch 200  
accuracy on train data: 0.7962  
learning rate: 0.0007071067811865475  
softmax loss: 1.7943773712320744e-06

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.8045  
learning rate: 0.0007053456158585983  
softmax loss: 2.1482919875104044e-06

accuracy on train data: 0.8045  
metrics for train data:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	1966
1	0.87	0.93	0.90	1845
2	0.66	0.80	0.72	1624
3	0.59	0.61	0.60	1932
4	0.91	0.72	0.80	2633
micro avg	0.80	0.80	0.80	10000
macro avg	0.80	0.81	0.80	10000
weighted avg	0.82	0.80	0.81	10000
samples avg	0.80	0.80	0.80	10000

accuracy on test data: 0.762  
metrics for test data:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	226
1	0.80	0.88	0.84	179
2	0.57	0.70	0.63	164
3	0.57	0.54	0.55	199
4	0.85	0.69	0.76	232
micro avg	0.76	0.76	0.76	1000
macro avg	0.75	0.76	0.75	1000
weighted avg	0.77	0.76	0.76	1000
samples avg	0.76	0.76	0.76	1000

Hidden layer: [512, 256, 128, 64]  
epoch 50  
accuracy on train data: 0.7216  
learning rate: 0.001414213562373095  
softmax loss: 1.9969842319606146e-05

epoch 100  
accuracy on train data: 0.7709  
learning rate: 0.001  
softmax loss: 1.154033885296016e-05

epoch 150  
accuracy on train data: 0.8061  
learning rate: 0.0008164965809277261  
softmax loss: 9.192962043513032e-07

epoch 200  
accuracy on train data: 0.8028  
learning rate: 0.0007071067811865475  
softmax loss: 4.717764805687347e-06

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.832  
learning rate: 0.0007053456158585983  
softmax loss: 9.181468735261794e-07

accuracy on train data: 0.832  
metrics for train data:

	precision	recall	f1-score	support
0	0.99	0.98	0.98	1988
1	0.93	0.91	0.92	2034
2	0.79	0.80	0.80	1932
3	0.57	0.71	0.63	1602
4	0.88	0.76	0.81	2444
micro avg	0.83	0.83	0.83	10000
macro avg	0.83	0.83	0.83	10000
weighted avg	0.85	0.83	0.84	10000
samples avg	0.83	0.83	0.83	10000

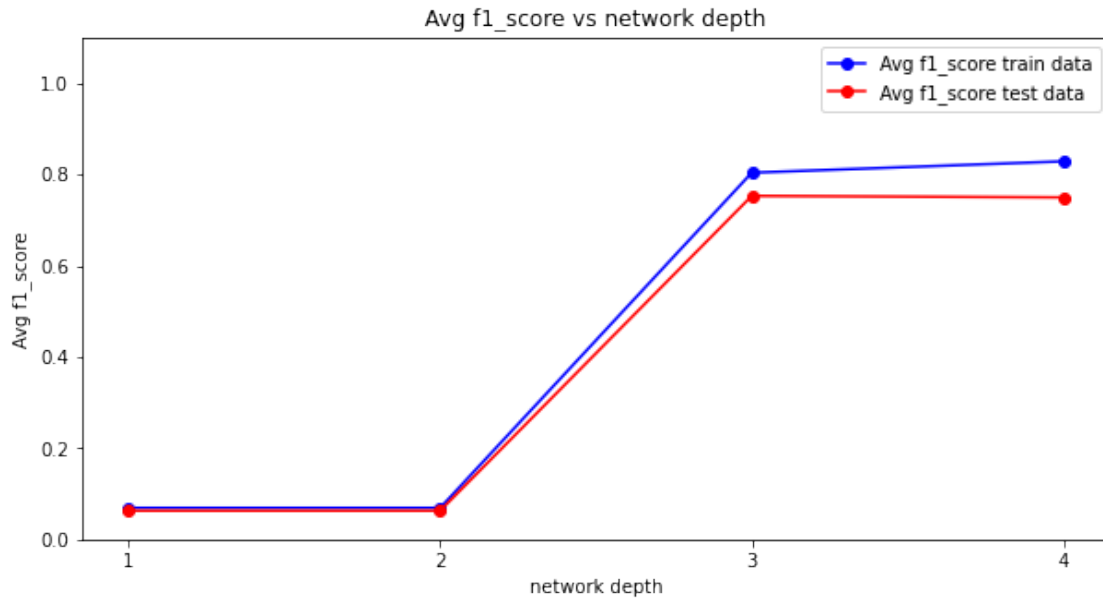
accuracy on test data: 0.761

metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.98	0.97	225
1	0.86	0.81	0.83	210
2	0.62	0.68	0.65	182
3	0.52	0.56	0.54	174
4	0.80	0.71	0.75	209
micro avg	0.76	0.76	0.76	1000
macro avg	0.75	0.75	0.75	1000
weighted avg	0.77	0.76	0.76	1000
samples avg	0.76	0.76	0.76	1000

```
[49]: #Plot avg f1_scores for different depth
plt.figure(figsize=(10,5))
plt.plot(network_depth, f1_score_train, marker='o', markersize=6, color='blue',
         ↳label='Avg f1_score train data')
plt.plot(network_depth, f1_score_test, marker='o', markersize=6, color='red',
         ↳label='Avg f1_score test data')

plt.title('Avg f1_score vs network depth')
plt.xlabel('network depth')
plt.ylabel('Avg f1_score')
plt.xticks(network_depth)
plt.ylim(0,1.1)
plt.legend()
plt.show()
```



### Relu activation function (with learning rate = 0.001)

```
[6]: hidden_layers = [[512], [512,256], [512,256,128], [512,256,128,64]]
network_depth = [1,2,3,4]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    NN = NeuralNetwork(n = 1024, n_hidden_nodes = hidden_layer , r = 5, M = 32)
    NN.train(X_train, y_train_onehot, epoch_mode= False, activation="relu",
    ↪adaptive_learning = True, alpha = 0.001, stopping_threshold = 5.0e-06,
    ↪printafter=50)
    y_pred_train, _ = NN.predict(X_train, activation="relu")
    y_pred_test, _ = NN.predict(X_test, activation="relu")

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
    get_metric(y_test_onehot, y_pred_test)
    f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))
```

```
print("\n")
```

Hidden layer: [512]

epoch 50

accuracy on train data: 0.6984

learning rate: 0.0001414213562373095

softmax loss: 3.2206525277455114e-05

epoch 100

accuracy on train data: 0.7177

learning rate: 0.0001

softmax loss: 2.8674372883269804e-05

epoch 150

accuracy on train data: 0.7355

learning rate: 8.164965809277261e-05

softmax loss: 2.6139527954960384e-05

epoch 200

accuracy on train data: 0.7488

learning rate: 7.071067811865475e-05

softmax loss: 2.3554841256167726e-05

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.7493

learning rate: 7.053456158585983e-05

softmax loss: 2.3523313554780805e-05

accuracy on train data: 0.7493

metrics for train data:

	precision	recall	f1-score	support
0	0.95	0.91	0.93	2050
1	0.78	0.77	0.78	1990
2	0.64	0.66	0.65	1904
3	0.54	0.62	0.58	1748
4	0.84	0.76	0.80	2308
micro avg	0.75	0.75	0.75	10000
macro avg	0.75	0.74	0.75	10000
weighted avg	0.76	0.75	0.75	10000

samples avg	0.75	0.75	0.75	10000
-------------	------	------	------	-------

accuracy on test data: 0.734

metrics for test data:

	precision	recall	f1-score	support
0	0.94	0.95	0.95	228
1	0.77	0.76	0.76	200
2	0.60	0.63	0.61	189
3	0.57	0.56	0.56	189
4	0.75	0.73	0.74	194

micro avg	0.73	0.73	0.73	1000
macro avg	0.73	0.72	0.73	1000
weighted avg	0.74	0.73	0.73	1000
samples avg	0.73	0.73	0.73	1000

Hidden layer: [512, 256]

epoch 50

accuracy on train data: 0.7485  
learning rate: 0.0001414213562373095  
softmax loss: 2.3803159130456424e-05

epoch 100

accuracy on train data: 0.7995  
learning rate: 0.0001  
softmax loss: 1.7794089613993414e-05

epoch 150

accuracy on train data: 0.8238  
learning rate: 8.164965809277261e-05  
softmax loss: 1.1009632399852487e-05

epoch 200

accuracy on train data: 0.842  
learning rate: 7.071067811865475e-05  
softmax loss: 8.309603697956972e-06

Convergence criteria satisfied!

epoch 201

accuracy on train data: 0.8419  
learning rate: 7.053456158585983e-05

softmax loss: 8.155999959208848e-06

accuracy on train data: 0.8419

metrics for train data:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1956
1	0.88	0.91	0.90	1910
2	0.77	0.81	0.78	1857
3	0.69	0.71	0.70	1963
4	0.89	0.80	0.84	2314
micro avg	0.84	0.84	0.84	10000
macro avg	0.84	0.84	0.84	10000
weighted avg	0.84	0.84	0.84	10000
samples avg	0.84	0.84	0.84	10000

accuracy on test data: 0.803

metrics for test data:

	precision	recall	f1-score	support
0	0.97	1.00	0.98	224
1	0.85	0.86	0.86	196
2	0.69	0.75	0.72	182
3	0.64	0.62	0.63	191
4	0.83	0.75	0.79	207
micro avg	0.80	0.80	0.80	1000
macro avg	0.80	0.80	0.80	1000
weighted avg	0.80	0.80	0.80	1000
samples avg	0.80	0.80	0.80	1000

Hidden layer: [512, 256, 128]

epoch 50

accuracy on train data: 0.7774

learning rate: 0.0001414213562373095

softmax loss: 2.9403732478028548e-05

epoch 100

accuracy on train data: 0.8168

learning rate: 0.0001

softmax loss: 1.2094483311416297e-05

epoch 150  
accuracy on train data: 0.847  
learning rate: 8.164965809277261e-05  
softmax loss: 8.964043796841776e-06

epoch 200  
accuracy on train data: 0.8617  
learning rate: 7.071067811865475e-05  
softmax loss: 5.319551939104827e-06

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.8618  
learning rate: 7.053456158585983e-05  
softmax loss: 5.11851412273001e-06

accuracy on train data: 0.8618  
metrics for train data:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1922
1	0.88	0.93	0.91	1872
2	0.82	0.83	0.82	1927
3	0.71	0.76	0.74	1893
4	0.92	0.81	0.86	2386
micro avg	0.86	0.86	0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000
samples avg	0.86	0.86	0.86	10000

accuracy on test data: 0.816  
metrics for test data:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	224
1	0.83	0.89	0.86	186
2	0.70	0.76	0.73	182
3	0.72	0.64	0.67	211
4	0.83	0.79	0.81	197
micro avg	0.82	0.82	0.82	1000
macro avg	0.81	0.81	0.81	1000
weighted avg	0.81	0.82	0.81	1000
samples avg	0.82	0.82	0.82	1000



Hidden layer: [512, 256, 128, 64]  
epoch 50  
accuracy on train data: 0.7796  
learning rate: 0.0001414213562373095  
softmax loss: 2.1823853735152896e-05

epoch 100  
accuracy on train data: 0.8303  
learning rate: 0.0001  
softmax loss: 9.080570686102212e-06

epoch 150  
accuracy on train data: 0.8567  
learning rate: 8.164965809277261e-05  
softmax loss: 5.53839396842743e-06

epoch 200  
accuracy on train data: 0.8847  
learning rate: 7.071067811865475e-05  
softmax loss: 5.457597634399534e-06

Convergence criteria satisfied!  
epoch 201  
accuracy on train data: 0.8908  
learning rate: 7.053456158585983e-05  
softmax loss: 3.781473057701032e-06

accuracy on train data: 0.8908  
metrics for train data:

	precision	recall	f1-score	support
0	0.99	0.98	0.99	1990
1	0.93	0.95	0.94	1941
2	0.85	0.87	0.86	1903
3	0.74	0.81	0.78	1832
4	0.94	0.84	0.89	2334
micro avg	0.89	0.89	0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

samples avg	0.89	0.89	0.89	10000
-------------	------	------	------	-------

accuracy on test data: 0.82

metrics for test data:


	precision	recall	f1-score	support
0	1.00	0.98	0.99	232
1	0.87	0.91	0.89	191
2	0.71	0.80	0.76	177
3	0.67	0.63	0.65	201
4	0.81	0.76	0.78	199

micro avg	0.82	0.82	0.82	1000
macro avg	0.81	0.82	0.81	1000
weighted avg	0.82	0.82	0.82	1000
samples avg	0.82	0.82	0.82	1000

```
[7]: #Plot avg f1_scores for different depth
plt.figure(figsize=(10,5))
plt.plot(network_depth, f1_score_train, marker='o', markersize=6, color='blue',
        ↳label='Avg f1_score train data')
plt.plot(network_depth, f1_score_test, marker='o', markersize=6, color='red',
        ↳label='Avg f1_score test data')

plt.title('Avg f1_score vs network depth')
plt.xlabel('network depth')
plt.ylabel('Avg f1_score')
plt.xticks(network_depth)
plt.ylim(0,1.1)
plt.legend()
plt.show()
```



### Neural Networks using scikit learn

```
[19]: hidden_layers = [[512], [512,256], [512,256,128], [512,256,128,64]]
network_depth = [1,2,3,4]
f1_score_train = []
f1_score_test = []
for hidden_layer in hidden_layers:
    print(f"Hidden layer: {hidden_layer}")
    clf = MLPClassifier(activation="relu", solver="sgd", alpha = 0,
    ↪batch_size=32, hidden_layer_sizes=np.array(hidden_layer),
    ↪learning_rate="invscaling", tol=5e-6, n_iter_no_change=5, verbose=True,
    ↪learning_rate_init=0.01).fit(X_train, y_train_onehot)
    y_pred_train = clf.predict(X_train)
    y_pred_test = clf.predict(X_test)

    print("accuracy on train data: ",accuracy_score(y_train_onehot,
    ↪y_pred_train))
    print("metrics for train data: ")
    get_metric(y_train_onehot, y_pred_train)
    f1_score_train.append(f1_score(y_train_onehot, y_pred_train,
    ↪average="macro"))

    print("accuracy on test data: ",accuracy_score(y_test_onehot, y_pred_test))
    print("metrics for test data: ")
    get_metric(y_test_onehot, y_pred_test)
    f1_score_test.append(f1_score(y_test_onehot, y_pred_test, average="macro"))
```

```
print("\n")
```

```
Hidden layer: [512]
Iteration 1, loss = 1.82100347
Iteration 2, loss = 1.71391782
Iteration 3, loss = 1.52315981
Iteration 4, loss = 1.51474497
Iteration 5, loss = 1.51029548
Iteration 6, loss = 1.50750493
Iteration 7, loss = 1.50565037
Iteration 8, loss = 1.50367111
Iteration 9, loss = 1.50203666
Iteration 10, loss = 1.50045418
Iteration 11, loss = 1.49933494
Iteration 12, loss = 1.49815999
Iteration 13, loss = 1.49698195
Iteration 14, loss = 1.49620743
Iteration 15, loss = 1.49550628
Iteration 16, loss = 1.49454919
Iteration 17, loss = 1.49363666
Iteration 18, loss = 1.49279152
Iteration 19, loss = 1.49200207
Iteration 20, loss = 1.49143824
Iteration 21, loss = 1.49060290
Iteration 22, loss = 1.48995872
Iteration 23, loss = 1.48938653
Iteration 24, loss = 1.48872835
Iteration 25, loss = 1.48810307
Iteration 26, loss = 1.48745704
Iteration 27, loss = 1.48675391
Iteration 28, loss = 1.48601222
Iteration 29, loss = 1.48555051
Iteration 30, loss = 1.48489258
Iteration 31, loss = 1.48442929
Iteration 32, loss = 1.48394169
Iteration 33, loss = 1.48350804
Iteration 34, loss = 1.48292553
Iteration 35, loss = 1.48245257
Iteration 36, loss = 1.48207413
Iteration 37, loss = 1.48157765
Iteration 38, loss = 1.48119172
Iteration 39, loss = 1.48078643
Iteration 40, loss = 1.48028934
Iteration 41, loss = 1.47994385
Iteration 42, loss = 1.47956204
Iteration 43, loss = 1.47913680
Iteration 44, loss = 1.47881800
```

Iteration 45, loss = 1.47835719  
Iteration 46, loss = 1.47806096  
Iteration 47, loss = 1.47762217  
Iteration 48, loss = 1.47733798  
Iteration 49, loss = 1.47684794  
Iteration 50, loss = 1.47658435  
Iteration 51, loss = 1.47617737  
Iteration 52, loss = 1.47584291  
Iteration 53, loss = 1.47550371  
Iteration 54, loss = 1.47513914  
Iteration 55, loss = 1.47476276  
Iteration 56, loss = 1.47438837  
Iteration 57, loss = 1.47401940  
Iteration 58, loss = 1.47368000  
Iteration 59, loss = 1.47319879  
Iteration 60, loss = 1.47285898  
Iteration 61, loss = 1.47256885  
Iteration 62, loss = 1.47207112  
Iteration 63, loss = 1.47172472  
Iteration 64, loss = 1.47138364  
Iteration 65, loss = 1.47103137  
Iteration 66, loss = 1.47073728  
Iteration 67, loss = 1.47036312  
Iteration 68, loss = 1.47003571  
Iteration 69, loss = 1.46979185  
Iteration 70, loss = 1.46951244  
Iteration 71, loss = 1.46925222  
Iteration 72, loss = 1.46899571  
Iteration 73, loss = 1.46865432  
Iteration 74, loss = 1.46840627  
Iteration 75, loss = 1.46819481  
Iteration 76, loss = 1.46778235  
Iteration 77, loss = 1.46766729  
Iteration 78, loss = 1.46732930  
Iteration 79, loss = 1.46709043  
Iteration 80, loss = 1.46691234  
Iteration 81, loss = 1.46659407  
Iteration 82, loss = 1.46647751  
Iteration 83, loss = 1.46619232  
Iteration 84, loss = 1.46591034  
Iteration 85, loss = 1.46570945  
Iteration 86, loss = 1.46540203  
Iteration 87, loss = 1.46522251  
Iteration 88, loss = 1.46499896  
Iteration 89, loss = 1.46479823  
Iteration 90, loss = 1.46448519  
Iteration 91, loss = 1.46429983  
Iteration 92, loss = 1.46408573

Iteration 93, loss = 1.46382916  
Iteration 94, loss = 1.46363847  
Iteration 95, loss = 1.46335830  
Iteration 96, loss = 1.46317960  
Iteration 97, loss = 1.46292187  
Iteration 98, loss = 1.46274574  
Iteration 99, loss = 1.46254516  
Iteration 100, loss = 1.46228647  
Iteration 101, loss = 1.46209452  
Iteration 102, loss = 1.46192999  
Iteration 103, loss = 1.46172054  
Iteration 104, loss = 1.46138636  
Iteration 105, loss = 1.46131177  
Iteration 106, loss = 1.46108355  
Iteration 107, loss = 1.46087098  
Iteration 108, loss = 1.46064470  
Iteration 109, loss = 1.46041485  
Iteration 110, loss = 1.46027431  
Iteration 111, loss = 1.46008784  
Iteration 112, loss = 1.45988650  
Iteration 113, loss = 1.45972962  
Iteration 114, loss = 1.45952104  
Iteration 115, loss = 1.45931163  
Iteration 116, loss = 1.45909816  
Iteration 117, loss = 1.45895664  
Iteration 118, loss = 1.45877721  
Iteration 119, loss = 1.45861600  
Iteration 120, loss = 1.45835821  
Iteration 121, loss = 1.45818034  
Iteration 122, loss = 1.45801117  
Iteration 123, loss = 1.45784181  
Iteration 124, loss = 1.45764247  
Iteration 125, loss = 1.45748175  
Iteration 126, loss = 1.45728778  
Iteration 127, loss = 1.45711550  
Iteration 128, loss = 1.45688891  
Iteration 129, loss = 1.45679229  
Iteration 130, loss = 1.45655648  
Iteration 131, loss = 1.45639470  
Iteration 132, loss = 1.45618799  
Iteration 133, loss = 1.45606132  
Iteration 134, loss = 1.45590004  
Iteration 135, loss = 1.45568361  
Iteration 136, loss = 1.45551304  
Iteration 137, loss = 1.45533163  
Iteration 138, loss = 1.45524969  
Iteration 139, loss = 1.45502739  
Iteration 140, loss = 1.45485948

Iteration 141, loss = 1.45471283  
Iteration 142, loss = 1.45452669  
Iteration 143, loss = 1.45431788  
Iteration 144, loss = 1.45420511  
Iteration 145, loss = 1.45401596  
Iteration 146, loss = 1.45389316  
Iteration 147, loss = 1.45370787  
Iteration 148, loss = 1.45356289  
Iteration 149, loss = 1.45343819  
Iteration 150, loss = 1.45324188  
Iteration 151, loss = 1.45308210  
Iteration 152, loss = 1.45292444  
Iteration 153, loss = 1.45276024  
Iteration 154, loss = 1.45260947  
Iteration 155, loss = 1.45243541  
Iteration 156, loss = 1.45227674  
Iteration 157, loss = 1.45219223  
Iteration 158, loss = 1.45195802  
Iteration 159, loss = 1.45185917  
Iteration 160, loss = 1.45166933  
Iteration 161, loss = 1.45155882  
Iteration 162, loss = 1.45138742  
Iteration 163, loss = 1.45124925  
Iteration 164, loss = 1.45108788  
Iteration 165, loss = 1.45094660  
Iteration 166, loss = 1.45077949  
Iteration 167, loss = 1.45066910  
Iteration 168, loss = 1.45050262  
Iteration 169, loss = 1.45032493  
Iteration 170, loss = 1.45023156  
Iteration 171, loss = 1.45006452  
Iteration 172, loss = 1.44994194  
Iteration 173, loss = 1.44978832  
Iteration 174, loss = 1.44964848  
Iteration 175, loss = 1.44951194  
Iteration 176, loss = 1.44936269  
Iteration 177, loss = 1.44922681  
Iteration 178, loss = 1.44907058  
Iteration 179, loss = 1.44894591  
Iteration 180, loss = 1.44877779  
Iteration 181, loss = 1.44865781  
Iteration 182, loss = 1.44852393  
Iteration 183, loss = 1.44839336  
Iteration 184, loss = 1.44824172  
Iteration 185, loss = 1.44808943  
Iteration 186, loss = 1.44795043  
Iteration 187, loss = 1.44782687  
Iteration 188, loss = 1.44768827

```

Iteration 189, loss = 1.44756264
Iteration 190, loss = 1.44745318
Iteration 191, loss = 1.44731446
Iteration 192, loss = 1.44718501
Iteration 193, loss = 1.44703131
Iteration 194, loss = 1.44688879
Iteration 195, loss = 1.44674095
Iteration 196, loss = 1.44659639
Iteration 197, loss = 1.44651981
Iteration 198, loss = 1.44630971
Iteration 199, loss = 1.44626250
Iteration 200, loss = 1.44608502

```

```

/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:691:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.

```

```
warnings.warn(
```

```
accuracy on train data: 0.3589
```

```
metrics for train data:
```

	precision	recall	f1-score	support
0	0.88	0.87	0.88	1990
1	0.21	0.75	0.33	557
2	0.03	0.58	0.07	117
3	0.00	0.00	0.00	0
4	0.66	0.73	0.69	1890
micro avg	0.36	0.79	0.49	4554
macro avg	0.36	0.59	0.39	4554
weighted avg	0.68	0.79	0.71	4554
samples avg	0.36	0.36	0.36	4554

```
accuracy on test data: 0.363
```

```
metrics for test data:
```

	precision	recall	f1-score	support
0	0.88	0.89	0.89	226
1	0.20	0.71	0.31	56
2	0.05	0.60	0.08	15
3	0.00	0.00	0.00	0
4	0.60	0.69	0.65	163
micro avg	0.36	0.79	0.50	460
macro avg	0.35	0.58	0.39	460
weighted avg	0.67	0.79	0.71	460
samples avg	0.36	0.36	0.36	460



Hidden layer: [512, 256]

```
/home/tkarthikeyan/.local/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in labels with no true samples.  
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/tkarthikeyan/.local/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in samples with no true labels.  
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/tkarthikeyan/.local/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in labels with no true samples.  
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/tkarthikeyan/.local/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in samples with no true labels.  
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
Iteration 1, loss = 1.75599291  
Iteration 2, loss = 1.50284355  
Iteration 3, loss = 1.39129507  
Iteration 4, loss = 1.38382212  
Iteration 5, loss = 1.37945388  
Iteration 6, loss = 1.37674041  
Iteration 7, loss = 1.37485859  
Iteration 8, loss = 1.37307523  
Iteration 9, loss = 1.37166307  
Iteration 10, loss = 1.37024469  
Iteration 11, loss = 1.36924451  
Iteration 12, loss = 1.36834121  
Iteration 13, loss = 1.36742281  
Iteration 14, loss = 1.36596267  
Iteration 15, loss = 1.36591150  
Iteration 16, loss = 1.36515085  
Iteration 17, loss = 1.36496215  
Iteration 18, loss = 1.36400063  
Iteration 19, loss = 1.36335653  
Iteration 20, loss = 1.36291710  
Iteration 21, loss = 1.36243440  
Iteration 22, loss = 1.36167984  
Iteration 23, loss = 1.36167079  
Iteration 24, loss = 1.36108227
```

Iteration 25, loss = 1.36048485  
Iteration 26, loss = 1.35995516  
Iteration 27, loss = 1.35961524  
Iteration 28, loss = 1.35919002  
Iteration 29, loss = 1.35920701  
Iteration 30, loss = 1.35863334  
Iteration 31, loss = 1.35818692  
Iteration 32, loss = 1.35783203  
Iteration 33, loss = 1.35752371  
Iteration 34, loss = 1.35744524  
Iteration 35, loss = 1.35693297  
Iteration 36, loss = 1.35671225  
Iteration 37, loss = 1.35632659  
Iteration 38, loss = 1.35607923  
Iteration 39, loss = 1.35565661  
Iteration 40, loss = 1.35532491  
Iteration 41, loss = 1.35524164  
Iteration 42, loss = 1.35503213  
Iteration 43, loss = 1.35473587  
Iteration 44, loss = 1.35456664  
Iteration 45, loss = 1.35412737  
Iteration 46, loss = 1.35385597  
Iteration 47, loss = 1.35384404  
Iteration 48, loss = 1.35366097  
Iteration 49, loss = 1.35333122  
Iteration 50, loss = 1.35296687  
Iteration 51, loss = 1.35278991  
Iteration 52, loss = 1.35252206  
Iteration 53, loss = 1.35240748  
Iteration 54, loss = 1.35215867  
Iteration 55, loss = 1.35190118  
Iteration 56, loss = 1.35157876  
Iteration 57, loss = 1.35146040  
Iteration 58, loss = 1.35123069  
Iteration 59, loss = 1.35098145  
Iteration 60, loss = 1.35093349  
Iteration 61, loss = 1.35068915  
Iteration 62, loss = 1.35031236  
Iteration 63, loss = 1.35013024  
Iteration 64, loss = 1.35012739  
Iteration 65, loss = 1.35003691  
Iteration 66, loss = 1.34982555  
Iteration 67, loss = 1.34947631  
Iteration 68, loss = 1.34929738  
Iteration 69, loss = 1.34929669  
Iteration 70, loss = 1.34913114  
Iteration 71, loss = 1.34883042  
Iteration 72, loss = 1.34868134

Iteration 73, loss = 1.34864712  
Iteration 74, loss = 1.34836011  
Iteration 75, loss = 1.34807917  
Iteration 76, loss = 1.34797559  
Iteration 77, loss = 1.34787465  
Iteration 78, loss = 1.34777258  
Iteration 79, loss = 1.34758600  
Iteration 80, loss = 1.34750898  
Iteration 81, loss = 1.34723970  
Iteration 82, loss = 1.34714625  
Iteration 83, loss = 1.34700567  
Iteration 84, loss = 1.34683882  
Iteration 85, loss = 1.34662978  
Iteration 86, loss = 1.34655793  
Iteration 87, loss = 1.34639897  
Iteration 88, loss = 1.34631632  
Iteration 89, loss = 1.34605786  
Iteration 90, loss = 1.34599202  
Iteration 91, loss = 1.34590211  
Iteration 92, loss = 1.34566993  
Iteration 93, loss = 1.34562009  
Iteration 94, loss = 1.34530316  
Iteration 95, loss = 1.34535392  
Iteration 96, loss = 1.34511840  
Iteration 97, loss = 1.34508257  
Iteration 98, loss = 1.34492088  
Iteration 99, loss = 1.34477977  
Iteration 100, loss = 1.34469043  
Iteration 101, loss = 1.34450016  
Iteration 102, loss = 1.34451317  
Iteration 103, loss = 1.34436293  
Iteration 104, loss = 1.34416883  
Iteration 105, loss = 1.34412784  
Iteration 106, loss = 1.34393020  
Iteration 107, loss = 1.34386049  
Iteration 108, loss = 1.34368240  
Iteration 109, loss = 1.34351548  
Iteration 110, loss = 1.34346316  
Iteration 111, loss = 1.34343423  
Iteration 112, loss = 1.34330118  
Iteration 113, loss = 1.34307329  
Iteration 114, loss = 1.34295323  
Iteration 115, loss = 1.34266210  
Iteration 116, loss = 1.34288639  
Iteration 117, loss = 1.34263167  
Iteration 118, loss = 1.34253032  
Iteration 119, loss = 1.34240821  
Iteration 120, loss = 1.34237846

Iteration 121, loss = 1.34209611  
Iteration 122, loss = 1.34211200  
Iteration 123, loss = 1.34199165  
Iteration 124, loss = 1.34185707  
Iteration 125, loss = 1.34172529  
Iteration 126, loss = 1.34168345  
Iteration 127, loss = 1.34150643  
Iteration 128, loss = 1.34146673  
Iteration 129, loss = 1.34139252  
Iteration 130, loss = 1.34123276  
Iteration 131, loss = 1.34114062  
Iteration 132, loss = 1.34104560  
Iteration 133, loss = 1.34099703  
Iteration 134, loss = 1.34089745  
Iteration 135, loss = 1.34064735  
Iteration 136, loss = 1.34067901  
Iteration 137, loss = 1.34052816  
Iteration 138, loss = 1.34047595  
Iteration 139, loss = 1.34022566  
Iteration 140, loss = 1.34035068  
Iteration 141, loss = 1.34006524  
Iteration 142, loss = 1.34015107  
Iteration 143, loss = 1.34002622  
Iteration 144, loss = 1.33980957  
Iteration 145, loss = 1.33978189  
Iteration 146, loss = 1.33974959  
Iteration 147, loss = 1.33956550  
Iteration 148, loss = 1.33959023  
Iteration 149, loss = 1.33943153  
Iteration 150, loss = 1.33934654  
Iteration 151, loss = 1.33923359  
Iteration 152, loss = 1.33923656  
Iteration 153, loss = 1.33907638  
Iteration 154, loss = 1.33897517  
Iteration 155, loss = 1.33882380  
Iteration 156, loss = 1.33879743  
Iteration 157, loss = 1.33873184  
Iteration 158, loss = 1.33852695  
Iteration 159, loss = 1.33854091  
Iteration 160, loss = 1.33844815  
Iteration 161, loss = 1.33841766  
Iteration 162, loss = 1.33822323  
Iteration 163, loss = 1.33813612  
Iteration 164, loss = 1.33808318  
Iteration 165, loss = 1.33807003  
Iteration 166, loss = 1.33790580  
Iteration 167, loss = 1.33785206  
Iteration 168, loss = 1.33775579

```

Iteration 169, loss = 1.33763462
Iteration 170, loss = 1.33757735
Iteration 171, loss = 1.33748661
Iteration 172, loss = 1.33754250
Iteration 173, loss = 1.33731174
Iteration 174, loss = 1.33731091
Iteration 175, loss = 1.33705106
Iteration 176, loss = 1.33725832
Iteration 177, loss = 1.33692821
Iteration 178, loss = 1.33708370
Iteration 179, loss = 1.33692998
Iteration 180, loss = 1.33672309
Iteration 181, loss = 1.33667763
Iteration 182, loss = 1.33663205
Iteration 183, loss = 1.33659152
Iteration 184, loss = 1.33656869
Iteration 185, loss = 1.33644886
Iteration 186, loss = 1.33631704
Iteration 187, loss = 1.33624661
Iteration 188, loss = 1.33621864
Iteration 189, loss = 1.33608505
Iteration 190, loss = 1.33600338
Iteration 191, loss = 1.33588587
Iteration 192, loss = 1.33582098
Iteration 193, loss = 1.33584034
Iteration 194, loss = 1.33574212
Iteration 195, loss = 1.33565059
Iteration 196, loss = 1.33556211
Iteration 197, loss = 1.33560856
Iteration 198, loss = 1.33546104
Iteration 199, loss = 1.33536024
Iteration 200, loss = 1.33527121

```

```

/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:691:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.

```

```
warnings.warn(
```

```
accuracy on train data: 0.5267
```

```
metrics for train data:
```

	precision	recall	f1-score	support
0	0.89	0.88	0.89	2008
1	0.64	0.70	0.67	1806
2	0.35	0.59	0.44	1168
3	0.06	0.56	0.11	227
4	0.68	0.73	0.70	1967

micro avg	0.53	0.73	0.61	7176
macro avg	0.53	0.69	0.56	7176
weighted avg	0.66	0.73	0.68	7176
samples avg	0.53	0.53	0.53	7176

accuracy on test data: 0.521

metrics for test data:

	precision	recall	f1-score	support
0	0.90	0.90	0.90	229
1	0.62	0.69	0.65	178
2	0.34	0.61	0.44	112
3	0.06	0.50	0.11	22
4	0.61	0.70	0.65	164

micro avg	0.52	0.74	0.61	705
macro avg	0.51	0.68	0.55	705
weighted avg	0.65	0.74	0.68	705
samples avg	0.52	0.52	0.52	705

Hidden layer: [512, 256, 128]

```
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in samples with no true labels.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in samples with no true labels.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
Iteration 1, loss = 1.76380936
Iteration 2, loss = 1.38822441
Iteration 3, loss = 1.37395561
Iteration 4, loss = 1.36949873
Iteration 5, loss = 1.36609622
Iteration 6, loss = 1.36384879
Iteration 7, loss = 1.36170042
Iteration 8, loss = 1.36104174
Iteration 9, loss = 1.35993561
Iteration 10, loss = 1.35869727
Iteration 11, loss = 1.35788140
Iteration 12, loss = 1.35715549
Iteration 13, loss = 1.35655835
Iteration 14, loss = 1.35555626
```

Iteration 15, loss = 1.35546035  
Iteration 16, loss = 1.35481342  
Iteration 17, loss = 1.35427231  
Iteration 18, loss = 1.35355794  
Iteration 19, loss = 1.35328877  
Iteration 20, loss = 1.35279381  
Iteration 21, loss = 1.35206872  
Iteration 22, loss = 1.35187235  
Iteration 23, loss = 1.35125906  
Iteration 24, loss = 1.35109893  
Iteration 25, loss = 1.35066431  
Iteration 26, loss = 1.35031724  
Iteration 27, loss = 1.34988212  
Iteration 28, loss = 1.34967275  
Iteration 29, loss = 1.34915685  
Iteration 30, loss = 1.34871293  
Iteration 31, loss = 1.34865636  
Iteration 32, loss = 1.34782173  
Iteration 33, loss = 1.34799764  
Iteration 34, loss = 1.34749533  
Iteration 35, loss = 1.34747387  
Iteration 36, loss = 1.34707925  
Iteration 37, loss = 1.34701910  
Iteration 38, loss = 1.34669471  
Iteration 39, loss = 1.34637387  
Iteration 40, loss = 1.34622964  
Iteration 41, loss = 1.34576911  
Iteration 42, loss = 1.34561575  
Iteration 43, loss = 1.34543691  
Iteration 44, loss = 1.34494329  
Iteration 45, loss = 1.34484796  
Iteration 46, loss = 1.34442891  
Iteration 47, loss = 1.34450047  
Iteration 48, loss = 1.34425555  
Iteration 49, loss = 1.34391610  
Iteration 50, loss = 1.34371431  
Iteration 51, loss = 1.34373249  
Iteration 52, loss = 1.34320654  
Iteration 53, loss = 1.34325544  
Iteration 54, loss = 1.34297260  
Iteration 55, loss = 1.34286864  
Iteration 56, loss = 1.34254717  
Iteration 57, loss = 1.34214021  
Iteration 58, loss = 1.34238252  
Iteration 59, loss = 1.34202606  
Iteration 60, loss = 1.34194018  
Iteration 61, loss = 1.34165714  
Iteration 62, loss = 1.34165045

Iteration 63, loss = 1.34128502  
Iteration 64, loss = 1.34113178  
Iteration 65, loss = 1.34106945  
Iteration 66, loss = 1.34079202  
Iteration 67, loss = 1.34061909  
Iteration 68, loss = 1.34038140  
Iteration 69, loss = 1.34046859  
Iteration 70, loss = 1.34008239  
Iteration 71, loss = 1.33991804  
Iteration 72, loss = 1.33994966  
Iteration 73, loss = 1.33973863  
Iteration 74, loss = 1.33955967  
Iteration 75, loss = 1.33927927  
Iteration 76, loss = 1.33915318  
Iteration 77, loss = 1.33905393  
Iteration 78, loss = 1.33881603  
Iteration 79, loss = 1.33875910  
Iteration 80, loss = 1.33869055  
Iteration 81, loss = 1.33849121  
Iteration 82, loss = 1.33841044  
Iteration 83, loss = 1.33829843  
Iteration 84, loss = 1.33780279  
Iteration 85, loss = 1.33802149  
Iteration 86, loss = 1.33781639  
Iteration 87, loss = 1.33781684  
Iteration 88, loss = 1.33736824  
Iteration 89, loss = 1.33746816  
Iteration 90, loss = 1.33722624  
Iteration 91, loss = 1.33701026  
Iteration 92, loss = 1.33701021  
Iteration 93, loss = 1.33679836  
Iteration 94, loss = 1.33670708  
Iteration 95, loss = 1.33646210  
Iteration 96, loss = 1.33652772  
Iteration 97, loss = 1.33630714  
Iteration 98, loss = 1.33620470  
Iteration 99, loss = 1.33623657  
Iteration 100, loss = 1.33598109  
Iteration 101, loss = 1.33585048  
Iteration 102, loss = 1.33569540  
Iteration 103, loss = 1.33551924  
Iteration 104, loss = 1.33556123  
Iteration 105, loss = 1.33533158  
Iteration 106, loss = 1.33514170  
Iteration 107, loss = 1.33500388  
Iteration 108, loss = 1.33476624  
Iteration 109, loss = 1.33489067  
Iteration 110, loss = 1.33468848



Iteration 111, loss = 1.33461397  
Iteration 112, loss = 1.33423274  
Iteration 113, loss = 1.33441666  
Iteration 114, loss = 1.33423589  
Iteration 115, loss = 1.33388573  
Iteration 116, loss = 1.33426413  
Iteration 117, loss = 1.33396289  
Iteration 118, loss = 1.33379433  
Iteration 119, loss = 1.33395570  
Iteration 120, loss = 1.33367703  
Iteration 121, loss = 1.33348642  
Iteration 122, loss = 1.33339904  
Iteration 123, loss = 1.33335125  
Iteration 124, loss = 1.33330120  
Iteration 125, loss = 1.33314045  
Iteration 126, loss = 1.33295453  
Iteration 127, loss = 1.33279385  
Iteration 128, loss = 1.33279488  
Iteration 129, loss = 1.33266766  
Iteration 130, loss = 1.33263197  
Iteration 131, loss = 1.33244764  
Iteration 132, loss = 1.33238185  
Iteration 133, loss = 1.33233309  
Iteration 134, loss = 1.33211635  
Iteration 135, loss = 1.33206129  
Iteration 136, loss = 1.33191412  
Iteration 137, loss = 1.33193500  
Iteration 138, loss = 1.33174531  
Iteration 139, loss = 1.33164275  
Iteration 140, loss = 1.33157977  
Iteration 141, loss = 1.33158093  
Iteration 142, loss = 1.33150074  
Iteration 143, loss = 1.33117412  
Iteration 144, loss = 1.33116357  
Iteration 145, loss = 1.33111280  
Iteration 146, loss = 1.33101557  
Iteration 147, loss = 1.33097566  
Iteration 148, loss = 1.33084151  
Iteration 149, loss = 1.33060484  
Iteration 150, loss = 1.33070459  
Iteration 151, loss = 1.33060298  
Iteration 152, loss = 1.33041688  
Iteration 153, loss = 1.33034815  
Iteration 154, loss = 1.33019880  
Iteration 155, loss = 1.33029594  
Iteration 156, loss = 1.33010516  
Iteration 157, loss = 1.32992872  
Iteration 158, loss = 1.32984731

```
Iteration 159, loss = 1.32972096
Iteration 160, loss = 1.32968075
Iteration 161, loss = 1.32959529
Iteration 162, loss = 1.32949004
Iteration 163, loss = 1.32923991
Iteration 164, loss = 1.32951352
Iteration 165, loss = 1.32934075
Iteration 166, loss = 1.32920078
Iteration 167, loss = 1.32911943
Iteration 168, loss = 1.32898543
Iteration 169, loss = 1.32894605
Iteration 170, loss = 1.32869095
Iteration 171, loss = 1.32877422
Iteration 172, loss = 1.32867062
Iteration 173, loss = 1.32861083
Iteration 174, loss = 1.32855360
Iteration 175, loss = 1.32839278
Iteration 176, loss = 1.32834345
Iteration 177, loss = 1.32824340
Iteration 178, loss = 1.32808126
Iteration 179, loss = 1.32813627
Iteration 180, loss = 1.32815307
Iteration 181, loss = 1.32796064
Iteration 182, loss = 1.32788895
Iteration 183, loss = 1.32779912
Iteration 184, loss = 1.32772729
Iteration 185, loss = 1.32754699
Iteration 186, loss = 1.32747568
Iteration 187, loss = 1.32727646
Iteration 188, loss = 1.32744426
Iteration 189, loss = 1.32735706
Iteration 190, loss = 1.32732263
Iteration 191, loss = 1.32716804
Iteration 192, loss = 1.32695542
Iteration 193, loss = 1.32709863
Iteration 194, loss = 1.32680616
Iteration 195, loss = 1.32680788
Iteration 196, loss = 1.32691740
Iteration 197, loss = 1.32686565
Iteration 198, loss = 1.32667633
Iteration 199, loss = 1.32653160
Iteration 200, loss = 1.32654795
```

```
/home/tkarthikeyan/.local/lib/python3.8/site-
```

```
packages/sklearn/neural_network/_multilayer_perceptron.py:691:
```

```
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
```

```
warnings.warn(
```

```

accuracy on train data: 0.5642
metrics for train data:
      precision    recall  f1-score   support

0         0.89        0.88        0.88        1985
1         0.67        0.70        0.68        1889
2         0.42        0.59        0.49        1384
3         0.14        0.54        0.23         531
4         0.71        0.71        0.71        2067


   micro avg       0.56       0.72       0.63       7856
   macro avg       0.56       0.68       0.60       7856
weighted avg       0.65       0.72       0.68       7856
 samples avg       0.56       0.56       0.56       7856

```

```

accuracy on test data: 0.562
metrics for test data:
      precision    recall  f1-score   support

0         0.90        0.91        0.90         225
1         0.64        0.68        0.66         184
2         0.40        0.61        0.48         130
3         0.18        0.54        0.27          63
4         0.63        0.68        0.66         173


   micro avg       0.56       0.73       0.63        775
   macro avg       0.55       0.69       0.59        775
weighted avg       0.63       0.73       0.67        775
 samples avg       0.56       0.56       0.56        775

```

Hidden layer: [512, 256, 128, 64]

```

/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in samples with no true labels.
Use `zero_division` parameter to control this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))
/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in samples with no true labels.
Use `zero_division` parameter to control this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))

```

```

Iteration 1, loss = 1.81368436
Iteration 2, loss = 1.44544842
Iteration 3, loss = 1.38917601
Iteration 4, loss = 1.38138721

```

Iteration 5, loss = 1.37689060  
Iteration 6, loss = 1.37331294  
Iteration 7, loss = 1.37047256  
Iteration 8, loss = 1.36867094  
Iteration 9, loss = 1.36701830  
Iteration 10, loss = 1.36629762  
Iteration 11, loss = 1.36428958  
Iteration 12, loss = 1.36411254  
Iteration 13, loss = 1.36277406  
Iteration 14, loss = 1.36201276  
Iteration 15, loss = 1.36122268  
Iteration 16, loss = 1.36026491  
Iteration 17, loss = 1.35992904  
Iteration 18, loss = 1.35911954  
Iteration 19, loss = 1.35855620  
Iteration 20, loss = 1.35809020  
Iteration 21, loss = 1.35724215  
Iteration 22, loss = 1.35677803  
Iteration 23, loss = 1.35652282  
Iteration 24, loss = 1.35581300  
Iteration 25, loss = 1.35539375  
Iteration 26, loss = 1.35481140  
Iteration 27, loss = 1.35403886  
Iteration 28, loss = 1.35396309  
Iteration 29, loss = 1.35331822  
Iteration 30, loss = 1.35319095  
Iteration 31, loss = 1.35287776  
Iteration 32, loss = 1.35243008  
Iteration 33, loss = 1.35148415  
Iteration 34, loss = 1.35169121  
Iteration 35, loss = 1.35138144  
Iteration 36, loss = 1.35108106  
Iteration 37, loss = 1.35077123  
Iteration 38, loss = 1.35049022  
Iteration 39, loss = 1.35019795  
Iteration 40, loss = 1.34961928  
Iteration 41, loss = 1.34928929  
Iteration 42, loss = 1.34920142  
Iteration 43, loss = 1.34881146  
Iteration 44, loss = 1.34820406  
Iteration 45, loss = 1.34864597  
Iteration 46, loss = 1.34793955  
Iteration 47, loss = 1.34757309  
Iteration 48, loss = 1.34729816  
Iteration 49, loss = 1.34697372  
Iteration 50, loss = 1.34694880  
Iteration 51, loss = 1.34656289  
Iteration 52, loss = 1.34625051

Iteration 53, loss = 1.34602164  
Iteration 54, loss = 1.34617614  
Iteration 55, loss = 1.34566605  
Iteration 56, loss = 1.34559899  
Iteration 57, loss = 1.34530104  
Iteration 58, loss = 1.34506928  
Iteration 59, loss = 1.34490077  
Iteration 60, loss = 1.34454417  
Iteration 61, loss = 1.34442064  
Iteration 62, loss = 1.34423867  
Iteration 63, loss = 1.34389345  
Iteration 64, loss = 1.34411251  
Iteration 65, loss = 1.34349359  
Iteration 66, loss = 1.34326719  
Iteration 67, loss = 1.34327519  
Iteration 68, loss = 1.34291595  
Iteration 69, loss = 1.34269664  
Iteration 70, loss = 1.34257381  
Iteration 71, loss = 1.34238822  
Iteration 72, loss = 1.34225701  
Iteration 73, loss = 1.34217715  
Iteration 74, loss = 1.34199308  
Iteration 75, loss = 1.34156418  
Iteration 76, loss = 1.34145023  
Iteration 77, loss = 1.34130128  
Iteration 78, loss = 1.34107888  
Iteration 79, loss = 1.34097266  
Iteration 80, loss = 1.34051020  
Iteration 81, loss = 1.34031797  
Iteration 82, loss = 1.34050280  
Iteration 83, loss = 1.34022977  
Iteration 84, loss = 1.34014982  
Iteration 85, loss = 1.33990893  
Iteration 86, loss = 1.33985747  
Iteration 87, loss = 1.33962806  
Iteration 88, loss = 1.33948541  
Iteration 89, loss = 1.33930273  
Iteration 90, loss = 1.33923741  
Iteration 91, loss = 1.33903170  
Iteration 92, loss = 1.33890641  
Iteration 93, loss = 1.33867310  
Iteration 94, loss = 1.33863891  
Iteration 95, loss = 1.33824019  
Iteration 96, loss = 1.33817394  
Iteration 97, loss = 1.33787108  
Iteration 98, loss = 1.33795980  
Iteration 99, loss = 1.33778427  
Iteration 100, loss = 1.33762673

Iteration 101, loss = 1.33748697  
Iteration 102, loss = 1.33726676  
Iteration 103, loss = 1.33727729  
Iteration 104, loss = 1.33722797  
Iteration 105, loss = 1.33697452  
Iteration 106, loss = 1.33682430  
Iteration 107, loss = 1.33675692  
Iteration 108, loss = 1.33645577  
Iteration 109, loss = 1.33642302  
Iteration 110, loss = 1.33632492  
Iteration 111, loss = 1.33604307  
Iteration 112, loss = 1.33606265  
Iteration 113, loss = 1.33586589  
Iteration 114, loss = 1.33588622  
Iteration 115, loss = 1.33560642  
Iteration 116, loss = 1.33555894  
Iteration 117, loss = 1.33544376  
Iteration 118, loss = 1.33528809  
Iteration 119, loss = 1.33515806  
Iteration 120, loss = 1.33508599  
Iteration 121, loss = 1.33489823  
Iteration 122, loss = 1.33480901  
Iteration 123, loss = 1.33486307  
Iteration 124, loss = 1.33455791  
Iteration 125, loss = 1.33450993  
Iteration 126, loss = 1.33431833  
Iteration 127, loss = 1.33432848  
Iteration 128, loss = 1.33407660  
Iteration 129, loss = 1.33384337  
Iteration 130, loss = 1.33399172  
Iteration 131, loss = 1.33385179  
Iteration 132, loss = 1.33338531  
Iteration 133, loss = 1.33348674  
Iteration 134, loss = 1.33337484  
Iteration 135, loss = 1.33334390  
Iteration 136, loss = 1.33319013  
Iteration 137, loss = 1.33311663  
Iteration 138, loss = 1.33289270  
Iteration 139, loss = 1.33284942  
Iteration 140, loss = 1.33269074  
Iteration 141, loss = 1.33260482  
Iteration 142, loss = 1.33249308  
Iteration 143, loss = 1.33238653  
Iteration 144, loss = 1.33241869  
Iteration 145, loss = 1.33235419  
Iteration 146, loss = 1.33195410  
Iteration 147, loss = 1.33200446  
Iteration 148, loss = 1.33184600

Iteration 149, loss = 1.33173865  
Iteration 150, loss = 1.33163852  
Iteration 151, loss = 1.33162311  
Iteration 152, loss = 1.33157170  
Iteration 153, loss = 1.33124986  
Iteration 154, loss = 1.33114563  
Iteration 155, loss = 1.33117362  
Iteration 156, loss = 1.33111677  
Iteration 157, loss = 1.33091936  
Iteration 158, loss = 1.33059599  
Iteration 159, loss = 1.33069087  
Iteration 160, loss = 1.33072471  
Iteration 161, loss = 1.33060597  
Iteration 162, loss = 1.33048843  
Iteration 163, loss = 1.33034758  
Iteration 164, loss = 1.33026873  
Iteration 165, loss = 1.32985650  
Iteration 166, loss = 1.33006810  
Iteration 167, loss = 1.33002381  
Iteration 168, loss = 1.33001853  
Iteration 169, loss = 1.32981868  
Iteration 170, loss = 1.32967690  
Iteration 171, loss = 1.32967191  
Iteration 172, loss = 1.32964922  
Iteration 173, loss = 1.32915260  
Iteration 174, loss = 1.32933370  
Iteration 175, loss = 1.32919457  
Iteration 176, loss = 1.32925131  
Iteration 177, loss = 1.32907655  
Iteration 178, loss = 1.32889414  
Iteration 179, loss = 1.32900540  
Iteration 180, loss = 1.32868908  
Iteration 181, loss = 1.32884859  
Iteration 182, loss = 1.32870995  
Iteration 183, loss = 1.32838215  
Iteration 184, loss = 1.32848429  
Iteration 185, loss = 1.32849773  
Iteration 186, loss = 1.32830025  
Iteration 187, loss = 1.32789618  
Iteration 188, loss = 1.32815089  
Iteration 189, loss = 1.32800030  
Iteration 190, loss = 1.32774316  
Iteration 191, loss = 1.32808793  
Iteration 192, loss = 1.32764538  
Iteration 193, loss = 1.32781256  
Iteration 194, loss = 1.32763628  
Iteration 195, loss = 1.32751946  
Iteration 196, loss = 1.32739331

```

Iteration 197, loss = 1.32729238
Iteration 198, loss = 1.32733673
Iteration 199, loss = 1.32719312
Iteration 200, loss = 1.32718437

/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:691:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
  warnings.warn(

```

```

accuracy on train data: 0.5846
metrics for train data:

```

	precision	recall	f1-score	support
0	0.89	0.88	0.88	1997
1	0.68	0.69	0.69	1958
2	0.46	0.58	0.51	1568
3	0.18	0.53	0.27	688
4	0.71	0.71	0.71	2064
micro avg	0.58	0.71	0.64	8275
macro avg	0.58	0.68	0.61	8275
weighted avg	0.66	0.71	0.67	8275
samples avg	0.58	0.58	0.58	8275

```

accuracy on test data: 0.581
metrics for test data:

```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	228
1	0.66	0.68	0.67	194
2	0.44	0.59	0.51	149
3	0.20	0.45	0.28	85
4	0.64	0.69	0.66	172
micro avg	0.58	0.70	0.64	828
macro avg	0.57	0.66	0.60	828
weighted avg	0.64	0.70	0.66	828
samples avg	0.58	0.58	0.58	828

```


/home/tkarthikeyan/.local/lib/python3.8/site-
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in samples with no true labels.
Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```



```
/home/tkarthikeyan/.local/lib/python3.8/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Recall  
and F-score are ill-defined and being set to 0.0 in samples with no true labels.  
Use `zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, msg_start, len(result))
```

```
[20]: #Plot avg f1_scores for different depth  
plt.figure(figsize=(10,5))  
plt.plot(network_depth, f1_score_train, marker='o', markersize=6, color='blue',  
        ↪label='Avg f1_score train data')  
plt.plot(network_depth, f1_score_test, marker='o', markersize=6, color='red',  
        ↪label='Avg f1_score test data')  
  
plt.title('Avg f1_score vs network depth')  
plt.xlabel('network depth')  
plt.ylabel('Avg f1_score')  
plt.xticks(network_depth)  
plt.ylim(0,1.1)  
plt.legend()  
plt.show()
```



```
[ ]:
```